

Mahendar Bathini

Hyderabad, Telangana | mahendargoud9866@gmail.com | 063041 60158 | linkedin.com/in/mahendar-bathini

Education

JNTU University, B.Tech in Computer Science Engineering – Hyderabad, Telangana
July 2019 – July 2023

Experience

AI Engineer, INNOVAPATH IT Solutions PVT Ltd – Hyderabad, Telangana

Jan 2025 – present

- Built a Retrieval-Augmented Generation (RAG) assistant to automate candidate shortlisting by matching resumes to job descriptions based on skills, experience, and location, reducing manual screening effort by ~40%.
- Designed and deployed a document ingestion pipeline using AWS (S3, Lambda, DLQs) and Docker to process and index 10K+ resumes and job descriptions with improved reliability and fault tolerance.
- Implemented a hybrid search system combining vector similarity (Sentence Transformers, Milvus) with keyword-based retrieval, increasing candidate-job match accuracy by 25% compared to keyword-only search.
- Developed LLM workflows using LangChain and AWS Bedrock for query refinement and response generation, enabling recruiters to retrieve ranked candidates with source-backed explanations, reducing turnaround time per search by 30%.

AI Intern, People Tech Group – Hyderabad, Telangana

Mar 2024 – Dec 2024

AI/ML and GenAI internship focused on backend development, RAG, LLM applications, and AWS.

- Assisted in developing REST APIs using FastAPI for ML model integration and backend services.
- Contributed to CI/CD workflows using GitHub Actions for testing and deployment automation.
- Helped implement Redis caching for frequently accessed endpoints to improve response times.
- Worked with Docker containers and assisted in deployments on Amazon Web Services Amazon EKS environments.

Projects

TalentScreen RAG system

RAG based AI assistant for Talent Management for recruiter.

- Built a Retrieval-Augmented Generation (RAG)-based AI assistant to support talent management and recruiter workflows.
- Developed a scalable data ingestion pipeline using AWS S3, AWS Lambda, Docker, and Dead Letter Queues (DLQs) for robust document processing.
- Leveraged LangChain and AWS Bedrock to orchestrate LLM workflows and enable context-aware query handling.
- Implemented document preprocessing and metadata extraction, generating embeddings using Sentence Transformers and storing them in a Milvus vector database.
- Designed a hybrid retrieval system combining semantic search (vector embeddings) and keyword-based methods to improve relevance.
- Integrated citation generation to provide traceable and explainable responses.
- Applied query optimization techniques to better interpret complex and ambiguous job descriptions.

TalentScreen Agentic RAG system (Phase 2)

Extended an existing Retrieval-Augmented Generation (RAG) system into a multi-agent architecture using LangGraph to support automated resume evaluation and candidate communication workflows.

- Designed and implemented multiple task-specific agents, including resume comparison, review insights generation, recommendation generation, and automated email response handling.
- Built an orchestration layer using a supervisor agent pattern to coordinate agent interactions, manage task flow, and handle fallback scenarios.
- Implemented short-term and long-term memory management for agents, enabling context retention across sessions and improving response consistency.

- Addressed prompt reliability and output consistency issues through iterative prompt tuning and structured response handling.

Skills

Programming Languages: Python, JavaScript

Frameworks: Pytorch, MLflow, Reactjs, Nodejs, Redux

Libraries: Langchain, Llamaindex, NLTK, HuggingFace, Transformers, Scikit-learn, NumPy, Pandas

LLMs: LLaMA, Anthropic Claude, BERT, Mistral, OpenAI

Database: SQL, PostgreSQL, MongoDB

AWS: EKS, Bedrock, RDS, Lambda, CloudWatch, S3 bucket, EC2, API Gateway